



Attention vs MLP Sublayer Decomposition of Activation Patching


Llama-3.3-70B




Mistral-Small-3.2-24B



ALLaM-7B



MedGemma-27B



● Attention-only patch ■ MLP-only patch ▲ Full layer (reference) - - - English baseline (100%) - - - Arabic baseline (0%)